

SALES FORECASTING USING MACHINE LEARNING

**A Machine Learning Approach for Predictive Sales Analytics and
Business Intelligence**

TECHNICAL PROJECT REPORT

Prepared By

Karishma

Bachelor of Technology

Artificial Intelligence and Data Science

Project Domain

Machine Learning • Data Analytics • Business Intelligence • Time Series Forecasting

Technology Stack

Python • Pandas • NumPy • Scikit-learn • Plotly • Streamlit • Jupyter Notebook • VS
Code

Version 1.0

2026

ACKNOWLEDGEMENT

I would like to express my sincere gratitude to everyone who contributed directly or indirectly to the successful completion of this project titled “**Sales Forecasting Using Machine Learning.**”

I am thankful to my teachers and mentors for their valuable guidance, encouragement, and continuous support throughout the development of this project. Their suggestions and constructive feedback helped me improve both my technical knowledge and problem-solving skills.

I would also like to thank my family and friends for their constant motivation, patience, and encouragement during every stage of this project.

Finally, I acknowledge the open-source Python community and the developers of libraries such as **Pandas, NumPy, Scikit-learn, Plotly, and Streamlit**, whose excellent documentation and tools greatly contributed to the successful implementation of this project.

ABSTRACT

Sales forecasting is a critical business process that helps organizations predict future sales based on historical data, enabling better planning, inventory management, budgeting, and strategic decision-making. Accurate forecasting allows businesses to optimize operations while reducing financial risks associated with overstocking or understocking.

This project presents a machine learning-based sales forecasting system developed using Python. The workflow includes data collection, preprocessing, exploratory data analysis (EDA), feature engineering, model training, performance evaluation, and future sales prediction. Historical sales data was analyzed to identify trends and seasonal patterns before applying machine learning algorithms to generate reliable forecasts.

To enhance usability, an interactive dashboard was developed using **Streamlit**. The dashboard provides key performance indicators (KPIs), historical sales trends, forecast visualizations, monthly analysis, and downloadable forecast results through an intuitive user interface.

The developed system demonstrates how machine learning and data visualization techniques can be effectively combined to support business intelligence and improve data-driven decision-making.

TABLE OF CONTENTS

SALES FORECASTING USING MACHINE LEARNING.....	1
A Machine Learning Approach for Predictive Sales Analytics and Business Intelligence.....	1
TECHNICAL PROJECT REPORT.....	1
Prepared By.....	1
Karishma.....	1
Project Domain.....	1
Technology Stack.....	1
Version 1.0.....	1
ACKNOWLEDGEMENT.....	2
ABSTRACT.....	2
TABLE OF CONTENTS.....	3
CHAPTER 1.....	5
PROJECT OVERVIEW.....	5
1.1 Introduction.....	5
1.2 Problem Statement.....	5
1.3 Objectives.....	6
1.4 Scope of the Project.....	6
1.5 Technologies Used.....	6
1.6 Project Workflow.....	7
CHAPTER 2.....	8
DATASET & METHODOLOGY.....	8
2.1 Dataset Overview.....	8
2.2 Dataset Features.....	8
2.3 Data Preprocessing.....	8
2.4 Feature Engineering.....	9
2.5 Methodology.....	9
CHAPTER 3.....	10
MODEL DEVELOPMENT.....	10
3.1 Machine Learning Models.....	10
3.2 Model Training.....	10
3.3 Model Evaluation.....	11
3.4 Forecast Generation.....	11
 Table.....	11

CHAPTER 4.....	12
DASHBOARD & BUSINESS INSIGHTS.....	12
4.1 Dashboard Overview.....	12
4.2 Dashboard Features.....	12
4.3 Business Insights.....	13
4.4 Dashboard Benefits.....	13
 CHAPTER 5.....	 14
CONCLUSION & FUTURE SCOPE.....	14
5.1 Conclusion.....	14
5.2 Future Scope.....	15
 References.....	 15

CHAPTER 1

PROJECT OVERVIEW

1.1 Introduction

Sales forecasting is an essential business process that estimates future sales using historical data. Accurate forecasting helps organizations improve inventory management, financial planning, demand prediction, and overall decision-making.

Machine learning techniques provide more reliable forecasting by identifying patterns and trends that traditional statistical methods may overlook. These techniques enable businesses to make informed decisions based on data-driven insights.

This project develops a machine learning-based sales forecasting system using Python. The project includes data preprocessing, exploratory data analysis, feature engineering, model development, evaluation, and future sales prediction. An interactive Streamlit dashboard has also been developed to visualize historical sales, forecast results, and business insights.

The proposed system demonstrates how machine learning can improve forecasting accuracy while providing an easy-to-use interface for business analysis.

1.2 Problem Statement

Businesses often face challenges in accurately predicting future sales due to changing customer demand, seasonal variations, and large volumes of historical data. Poor forecasting can lead to inventory shortages, excess stock, and inefficient financial planning.

Traditional forecasting methods may not capture hidden patterns within the data. Therefore, there is a need for a machine learning-based forecasting system capable of generating reliable predictions while providing meaningful visual insights through an interactive dashboard.

1.3 Objectives

The primary objectives of this project are:

- To analyze historical sales data for identifying business trends and patterns.
- To preprocess and clean the dataset for accurate analysis.
- To perform feature engineering for improving prediction accuracy.
- To build machine learning models for sales forecasting.

- To evaluate the performance of forecasting models using appropriate evaluation metrics.
- To visualize historical and forecasted sales through an interactive Streamlit dashboard.
- To provide meaningful business insights that support data-driven decision-making.

1.4 Scope of the Project

This project focuses on developing a machine learning-based sales forecasting system capable of predicting future sales using historical data. The system includes data preprocessing, feature engineering, model development, evaluation, and dashboard visualization.

Although the project uses a sample sales dataset, the methodology can be applied to retail, e-commerce, manufacturing, healthcare, and other business domains requiring sales forecasting.

1.5 Technologies Used

Technology	Purpose
Python	Programming Language
Pandas	Data Manipulation and Analysis
NumPy	Numerical Computation
Matplotlib	Data Visualization
Plotly	Interactive Charts
Scikit-learn	Machine Learning Models

Streamlit	Dashboard Development
Jupyter Notebook	Data Analysis & Model Development
VS Code	Code Editing
Git & GitHub	Version Control and Project Management

1.6 Project Workflow

The project follows a structured workflow to transform raw sales data into meaningful business insights and future sales predictions. Each stage contributes to improving the overall performance and reliability of the forecasting system.

The workflow consists of the following steps:

Data Collection



Data Preprocessing



Exploratory Data Analysis



Feature Engineering



Model Building



Model Evaluation



Sales Forecasting



Dashboard Development

CHAPTER 2

DATASET & METHODOLOGY

2.1 Dataset Overview

The sales forecasting model was developed using a historical sales dataset containing transactional sales records. The dataset consists of order-related information such as sales amount, order date, customer details, product categories, and geographical information. Historical sales data serves as the foundation for identifying trends, seasonal patterns, and business behaviour required for accurate forecasting.

The dataset was imported into Python and analyzed using the Pandas library before undergoing preprocessing and feature engineering.

2.2 Dataset Features

The dataset contains several important attributes that contribute to sales forecasting. These include order information, customer details, product categories, geographical information, and sales values. The **Order Date** field plays a significant role in extracting time-based information such as year, month, week, and day, while the **Sales** column serves as the target variable for prediction. Additional information related to products, regions, and customers provides useful context for understanding sales patterns and business trends.

2.3 Data Preprocessing

Before model development, the dataset was cleaned and prepared for analysis. The preprocessing stage included checking for missing values, removing duplicate records, converting date columns into the DateTime format, verifying data types, and sorting the dataset chronologically.

These preprocessing steps ensured data consistency and improved the quality of the dataset for feature engineering and machine learning.

2.4 Feature Engineering

Feature engineering was performed to improve the predictive capability of the forecasting models. Several time-based and statistical features were generated from the historical sales data.

The engineered features include:

- Year
- Quarter
- Month
- Week
- Day
- Day of Week
- Is Weekend
- Lag 1
- Lag 7
- Lag 30
- Rolling Mean (7 Days)
- Rolling Mean (30 Days)
- Rolling Standard Deviation
- Expanding Mean
- Month Sin
- Month Cos
- Day Sin
- Day Cos

These features help the machine learning models capture temporal trends, seasonality, and historical sales behaviour more effectively.

2.5 Methodology

The project follows a systematic workflow beginning with data collection and preprocessing. The cleaned dataset is analyzed through exploratory data analysis to identify important sales trends and patterns. Feature engineering techniques are applied to generate meaningful input variables for the forecasting models. Machine learning algorithms are then trained and evaluated before generating future sales forecasts. Finally, the results are presented through an interactive Streamlit dashboard for business analysis.

```
[4]: # Display the first 5 rows of the dataset
df.head()
```

	Order ID	Order Date	Ship Date	Customer Name	Country	City	State	Category	Product Name	Sales	Quantity	
0	CA-2013-138688	2013-06-13	2013-06-17	Darrin Van Huff	United States	Los Angeles	California	Labels	Self-Adhesive Address Labels for Typewriters b...	14.620	2	6
1	CA-2011-115812	2011-06-09	2011-06-14	Brosina Hoffman	United States	Los Angeles	California	Furnishings	Eldon Expressions Wood and Plastic Desk Access...	48.860	7	14
2	CA-2011-115812	2011-06-09	2011-06-14	Brosina Hoffman	United States	Los Angeles	California	Art	Newell 322	7.280	4	1
3	CA-2011-115812	2011-06-09	2011-06-14	Brosina Hoffman	United States	Los Angeles	California	Phones	Mitel 5320 IP Phone VoIP phone	907.152	4	90
4	CA-2011-115812	2011-06-09	2011-06-14	Brosina Hoffman	United States	Los Angeles	California	Binders	DXL Angle-View Binders with Locking Rings by S...	18.504	3	5

CHAPTER 3

MODEL DEVELOPMENT

3.1 Machine Learning Models

The cleaned and feature-engineered dataset was used to train machine learning models for predicting future sales. The dataset was divided into training and testing sets to evaluate the performance of each model. Different regression algorithms were implemented and compared to identify the most suitable model for sales forecasting.

The primary objective was to develop a forecasting model capable of accurately predicting future sales based on historical sales patterns and engineered features.

3.2 Model Training

Two supervised machine learning algorithms were implemented during this project:

- **Linear Regression**
- **Random Forest Regression**

Linear Regression was used as a baseline model because of its simplicity and ability to model linear relationships between variables. Random Forest Regression was selected due to its capability to capture complex, non-linear relationships and improve prediction accuracy through an ensemble learning approach.

Both models were trained using the processed dataset and evaluated on unseen test data to measure their forecasting performance.

3.3 Model Evaluation

The performance of the developed machine learning models was evaluated using standard regression metrics to determine their forecasting accuracy. Three evaluation metrics were considered:

- **Mean Absolute Error (MAE):** Measures the average absolute difference between the actual and predicted sales values.
- **Root Mean Square Error (RMSE):** Measures the magnitude of prediction errors while giving higher importance to larger errors.
- **Coefficient of Determination (R² Score):** Indicates how well the model explains the variation in the target variable.

The evaluation results of both models are presented in Table


3.4 Forecast Generation

The evaluation results indicate that Linear Regression achieved better performance compared to Random Forest Regression on the given dataset. It produced lower prediction errors, as reflected by its MAE and RMSE values, and achieved a higher R² Score, indicating a better fit to the historical sales data.

Based on these results, the Linear Regression model was selected for generating future sales forecasts. The predicted sales values were stored and integrated into the Streamlit dashboard, where users can visualize forecast trends, analyze business performance, and download the forecasted data for further analysis.

 **Table**

Model	MAE	RMSE	R ² Score
Linear Regression	778.454024	1234.072821	0.200447
Random Forest	992.659382	1377.013528	0.018862



CHAPTER 4

DASHBOARD & BUSINESS INSIGHTS

4.1 Dashboard Overview

An interactive dashboard was developed using **Streamlit** to present the forecasting results in a simple and user-friendly manner. The dashboard enables users to explore historical sales trends, view future sales predictions, analyze business performance, and download forecasted data for further use.

The dashboard is designed to provide decision-makers with meaningful insights through interactive visualizations and key performance indicators (KPIs). It allows users to understand historical sales behavior and evaluate future sales trends without requiring technical knowledge of machine learning algorithms.

4.2 Dashboard Features

The developed dashboard provides the following features:

- Displays key sales performance indicators.
- Visualizes historical sales trends.
- Presents future sales forecasts generated by the machine learning model.
- Shows monthly sales analysis through interactive charts.
- Displays forecasted sales data in tabular format.
- Allows users to download forecast results as a CSV file.
- Provides an intuitive and responsive user interface for business analysis.




Figure 4.1: Sales Forecasting Dashboard

4.3 Business Insights

Based on the historical sales analysis and forecasting results, several meaningful business insights were identified:

- Historical sales data exhibits clear fluctuations over time, indicating the presence of seasonal demand.
- Forecasted sales trends help estimate future business performance.
- Machine learning enables faster and more reliable sales prediction compared to manual forecasting methods.
- Interactive visualizations improve understanding of business trends and support data-driven decision-making.
- Forecast reports can assist organizations in inventory planning, budgeting, and resource allocation.

4.4 Dashboard Benefits

The developed dashboard offers several advantages for business users:

- Simplifies sales analysis through interactive visualizations.
- Improves accessibility of forecasting results.
- Supports strategic business planning.
- Reduces manual effort in report generation.
- Enhances decision-making using machine learning predictions.



Figure 4.2: Forecast Visualization and Business Insights

CHAPTER 5

CONCLUSION & FUTURE SCOPE

5.1 Conclusion

This project successfully developed a machine learning-based sales forecasting system capable of predicting future sales using historical sales data. The project followed a structured workflow that included data preprocessing, exploratory data analysis, feature engineering, model development, performance evaluation, and dashboard visualization.

Among the evaluated models, **Linear Regression** achieved better performance than **Random Forest Regression**, producing lower prediction errors and a higher R^2 Score. Based on these results, Linear Regression was selected as the final forecasting model.

An interactive Streamlit dashboard was developed to present historical sales trends, forecasted sales, and business insights in a user-friendly manner. The dashboard allows users to analyze sales performance, visualize forecasting results, and download prediction data for further analysis.

Overall, the project demonstrates the practical application of machine learning and business intelligence techniques in solving real-world sales forecasting problems and supporting data-driven decision-making.

5.2 Future Scope

Although the developed system provides reliable forecasting results, several improvements can further enhance its performance and usability:

- Implementation of advanced forecasting models such as XGBoost, LightGBM, and LSTM.
- Integration with real-time business databases for live forecasting.
- Deployment of the dashboard on cloud platforms for online accessibility.
- Addition of automated report generation and email notifications.
- Inclusion of external factors such as holidays, promotions, weather conditions, and market trends to improve forecasting accuracy.
- Development of role-based dashboards for different business users.

These enhancements would increase the scalability and practical applicability of the forecasting system in real-world business environments.

References

1. Pedregosa, F., et al. (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research.
2. McKinney, W. (2022). *Python for Data Analysis* (3rd Edition). O'Reilly Media.
3. Pandas Documentation – <https://pandas.pydata.org>
4. NumPy Documentation – <https://numpy.org>
5. Scikit-learn Documentation – <https://scikit-learn.org>
6. Streamlit Documentation – <https://streamlit.io>
7. Plotly Documentation – <https://plotly.com/python>
8. Kaggle Datasets – <https://www.kaggle.com>